

✓ Klasifikasi Model Dengan Logistic Regression Dan KNN pada Dataset Iris TASK 1

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Machine Learning Week 1 Task 1 - Azmi Taquiddin Syah - 1103213078

✓ Import libraries Yang Dibutuhkan

Library yang di butuhkan Sebagai Berikut

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
import scipy
import statsmodels.formula.api as smf
import statsmodels.api as sm
from sklearn import model_selection, preprocessing, feature_selection, ensemble, linear_model, metrics, decomposition
```

```
column_names = ['tahun'] + [f'x{i}' for i in range(1, 91)]
```

```
dataset = "/content/sample_data/RegresiUTSTelkom.csv"
data = pd.read_csv(dataset,names=column_names)
```

```
data.head()
```

↗

	tahun	x1	x2	x3	x4	x5	x6	x7	x8	x9	...	x81	x82	x83
0	2001	49.94357	21.47114	73.07750	8.74861	-17.40628	-13.09905	-25.01202	-12.23257	7.83089	...	13.01620	-54.40548	58.99367
1	2001	48.73215	18.42930	70.32679	12.94636	-10.32437	-24.83777	8.76630	-0.92019	18.76548	...	5.66812	-19.68073	33.04964
2	2001	50.95714	31.85602	55.81851	13.41693	-6.57898	-18.54940	-3.27872	-2.35035	16.07017	...	3.03800	26.05866	-50.92779
3	2001	48.24750	-1.89837	36.29772	2.58776	0.97170	-26.21683	5.05097	-10.34124	3.55005	...	34.57337	-171.70734	-16.96705
4	2001	50.97020	42.20998	67.09964	8.46791	-15.85279	-16.81409	-12.48207	-9.37636	12.63699	...	9.92661	-55.95724	64.92712

5 rows × 91 columns

```
data.describe()
```

↗

	tahun	x1	x2	x3	x4	x5	x6	x7
count	515345.000000	515345.000000	515345.000000	515345.000000	515345.000000	515345.000000	515345.000000	515345.000000
mean	1998.397082	43.387126	1.289554	8.658347	1.164124	-6.553601	-9.521975	-2.391089
std	10.931046	6.067558	51.580351	35.268585	16.322790	22.860785	12.857751	14.571873
min	1922.000000	1.749000	-337.092500	-301.005060	-154.183580	-181.953370	-81.794290	-188.214000
25%	1994.000000	39.954690	-26.059520	-11.462710	-8.487500	-20.666450	-18.440990	-10.780600
50%	2002.000000	44.258500	8.417850	10.476320	-0.652840	-6.007770	-11.188390	-2.046670
75%	2006.000000	47.833890	36.124010	29.764820	8.787540	7.741870	-2.388960	6.508580
max	2011.000000	61.970140	384.065730	322.851430	335.771820	262.068870	166.236890	172.402680

8 rows × 91 columns

```
data.tail()
```

	tahun	x1	x2	x3	x4	x5	x6	x7	x8	x9	...	x81	x82	
515340	2006	51.28467	45.88068	22.19582	-5.53319	-3.61835	-16.36914	2.12652	5.18160	-8.66890	...	4.81440	-3.75991	-3
515341	2006	49.87870	37.93125	18.65987	-3.63581	-27.75665	-18.52988	7.76108	3.56109	-2.50351	...	32.38589	-32.75535	-6
515342	2006	45.12852	12.65758	-38.72018	8.80882	-29.29985	-2.28706	-18.40424	-22.28726	-4.52429	...	-18.73598	-71.15954	-12
515343	2006	44.16614	32.38368	-3.34971	-2.49165	-19.59278	-18.67098	8.78428	4.02039	-12.01230	...	67.16763	282.77624	-
515344	2005	51.85726	59.11655	26.39436	-5.46030	-20.69012	-19.95528	-6.72771	2.29590	10.31018	...	-11.50511	-69.18291	6

5 rows × 91 columns

View NaN values

print("Rows with NaN values:")

print(data[data.isna().any(axis=1)])


View duplicates

print("\nDuplicate rows:")

print(data[data.duplicated()])

Drop duplicates

data = data.drop_duplicates()

 [0 rows x 91 columns]

Duplicate rows:

	tahun	x1	x2	x3	x4	x5	x6	
5551	1973	41.22353	3.20571	43.66712	7.81090	-27.93823	-2.67931	
9941	1999	46.98706	38.24010	22.51761	7.24891	-3.88296	-5.25372	
9942	1998	43.28314	25.37917	6.80249	16.41132	-14.76744	3.84884	
10071	2006	45.88913	3.50835	-16.79630	2.95203	-12.71814	-10.46804	
17330	2008	31.59176	-43.16626	-53.11768	-7.08228	36.23470	-25.61305	
...	
500568	1988	38.97271	35.64567	-57.35630	-12.21976	-46.33510	-2.40727	
505982	2002	41.71236	-74.26705	-1.79009	2.35874	-41.19781	4.17658	
507253	2005	45.80278	31.58951	-7.59075	-14.01903	-18.81706	-7.66548	
507479	2008	43.46480	-10.83611	31.90017	-13.44218	-11.36997	2.02308	
508200	1992	39.18044	-23.14181	6.16137	2.35418	-9.26563	-6.34051	
		x7	x8	x9	...	x81	x82	x83 \
5551	-18.24996	-10.16836	10.21118	...	20.80983	-46.42137	-33.63520	
9941	12.18594	0.18605	12.75414	...	10.74268	-50.30651	-7.37361	
9942	11.74961	2.87696	5.98809	...	8.24654	-76.93471	-130.44354	
10071	0.08869	16.29474	-0.49840	...	12.71373	-163.83204	-156.44085	
17330	-31.75173	-11.16459	22.02261	...	27.60406	-256.30083	-11.62091	
...	
500568	-13.88561	-1.84301	0.86373	...	11.23003	-129.69128	-66.70936	
505982	-13.86373	2.79708	-1.92353	...	-1.52568	-341.90923	77.86735	
507253	-9.95065	-2.58846	4.83188	...	63.15197	10.50641	247.87142	
507479	-27.88310	-4.38305	20.88566	...	52.10270	74.33104	128.10386	
508200	-3.67249	-24.84162	-2.77967	...	20.63775	-40.43440	-16.27189	
		x84	x85	x86	x87	x88	x89	\
5551	-14.65970	1.63648	52.43969	-17.95543	-2.48364	-104.09383		
9941	95.50208	5.34527	-22.54009	95.45694	1.57117	-121.46786		
9942	49.00444	5.17171	-22.77865	196.49517	9.08196	-81.65460		
10071	165.20790	1.67669	-64.35191	94.78478	-5.55255	23.30911		
17330	119.28601	17.95582	51.36605	-116.69768	4.80931	114.35652		
...		
500568	13.65987	-5.99570	139.55599	-227.12107	11.44736	-82.32164		
505982	208.76815	-19.98203	64.68677	-147.48197	14.85343	-76.69082		
507253	66.71739	12.37958	6.93733	104.76817	22.78140	-151.14191		
507479	35.25644	14.86020	164.93570	25.08485	9.54500	-111.64292		
508200	48.69949	10.15451	111.65179	-47.06901	1.20738	129.55055		
		x90						
5551	6.01573							
9941	-8.07735							
9942	-2.64752							
10071	-6.75154							
17330	-11.11994							
...	...							
500568	-5.94241							
505982	-16.05562							
507253	19.05184							
507479	16.17438							
508200	3.48504							

[214 rows x 91 columns]

```
print("\nData after dropping duplicates:")
print(data.tail())
```



```
Data after dropping duplicates:
      tahun      x1      x2      x3      x4      x5      x6  \
515340    2006  51.28467  45.88068  22.19582 -5.53319 -3.61835 -16.36914
515341    2006  49.87870  37.93125  18.65987 -3.63581 -27.75665 -18.52988
515342    2006  45.12852  12.65758 -38.72018  8.80882 -29.29985 -2.28706
515343    2006  44.16614  32.38368 -3.34971 -2.49165 -19.59278 -18.67098
515344    2005  51.85726  59.11655  26.39436 -5.46030 -20.69012 -19.95528

      x7      x8      x9  ...      x81      x82      x83  \
515340  2.12652  5.18160 -8.66890  ...  4.81440 -3.75991 -30.92584
515341  7.76108  3.56109 -2.50351  ...  32.38589 -32.75535 -61.05473
515342 -18.40424 -22.28726 -4.52429  ... -18.73598 -71.15954 -123.98443
515343  8.78428  4.02039 -12.01230  ...  67.16763  282.77624 -4.63677
515344 -6.72771  2.29590  10.31018  ... -11.50511 -69.18291  60.58456

      x84      x85      x86      x87      x88      x89      x90
515340  26.33968 -5.03390  21.86037 -142.29410  3.42901 -41.14721 -15.46052
515341  56.65182 15.29965  95.88193 -10.63242 12.96552  92.11633  10.88815
515342 121.26989 10.89629  34.62409 -248.61020 -6.07171  53.96319 -8.09364
515343 144.00125 21.62652 -29.72432  71.47198  20.32240  14.83107  39.74909
515344  28.64599 -4.39620 -64.56491 -45.61012 -5.51512  32.35602 12.17352

[5 rows x 91 columns]
```

EDA

```
fig, ax = plt.subplots(figsize=(15,8))

ax.set_title('Distribution of songs per release tahun', fontsize=15)
variable = data['tahun']

sns.distplot(variable, hist=True, kde=True, kde_kws={"shade": True}, ax=ax)
des = data['tahun'].describe()
ax.axvline(des["25%"], ls='--')
ax.axvline(des["mean"], ls='--')
ax.axvline(des["75%"], ls='--')
ax.grid(True)

des = round(des).apply(lambda x: str(x))
box = '\n'.join(("min: "+des["min"], "25%: "+des["25%"], "mean: "+des["mean"], "75%: "+des["75%"], "max: "+des["max"]))
ax.text(0.20, 0.95, box, transform=ax.transAxes, fontsize=10, va='top', ha="right", bbox=dict(boxstyle='round', facecolor='white', alpha=0.5))
```



<ipython-input-10-05c37db356bd>:6: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

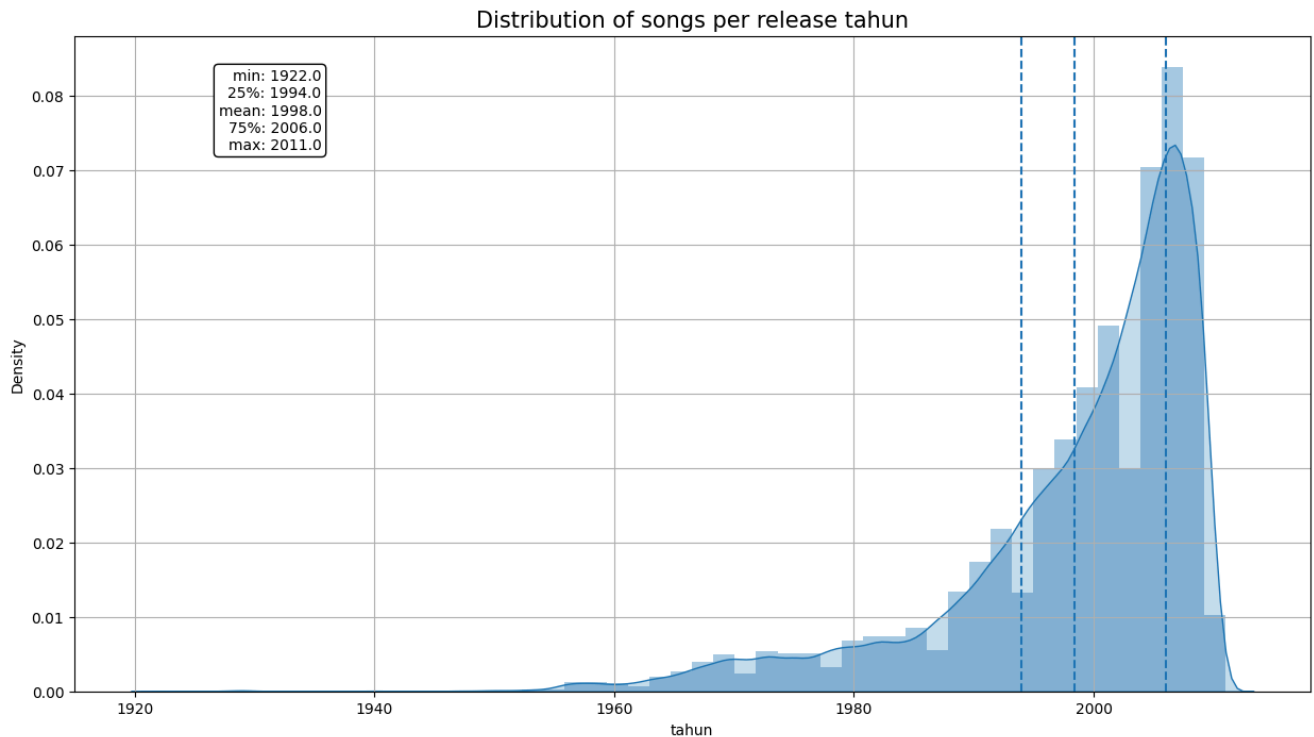
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(variable, hist=True, kde=True, kde_kws={"shade": True}, ax=ax)
/usr/local/lib/python3.10/dist-packages/seaborn/distributions.py:2496: FutureWarning:
```

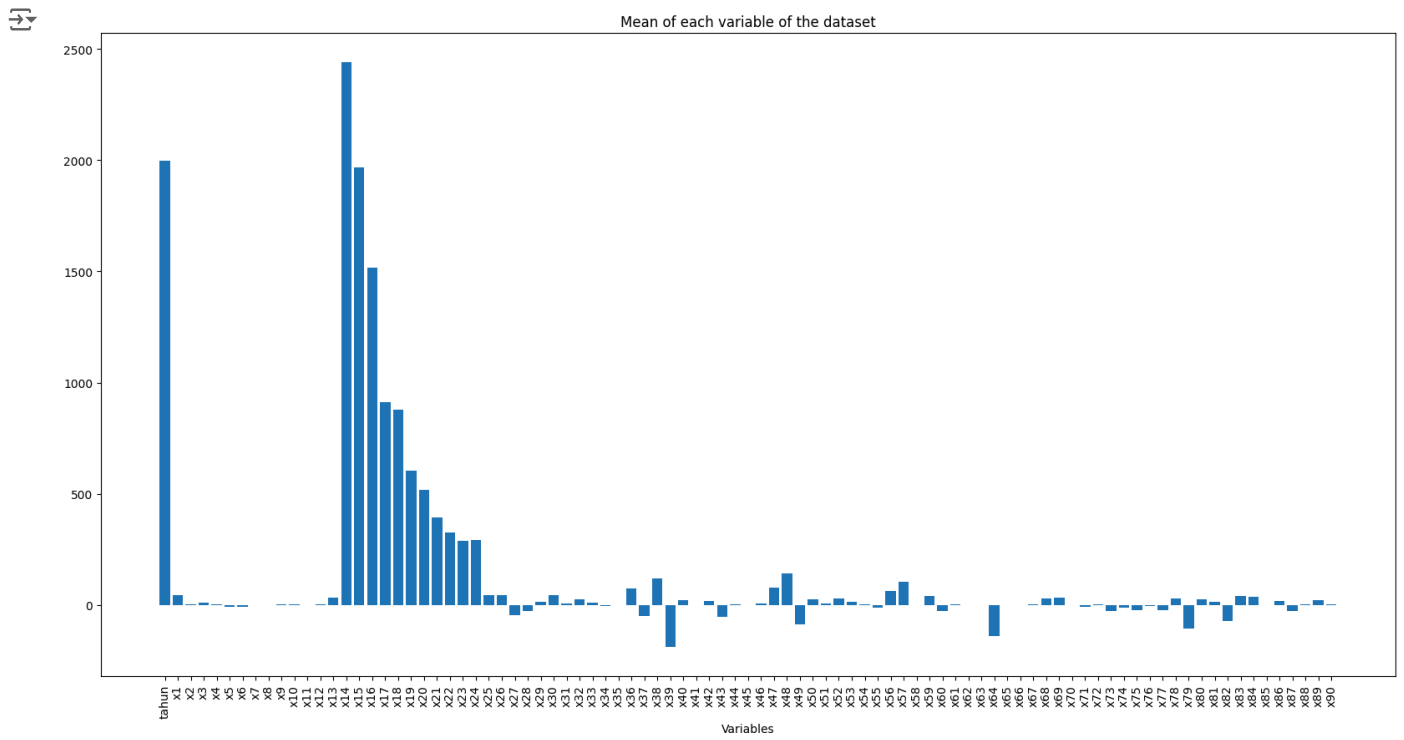
`shade` is now deprecated in favor of `fill`; setting `fill=True`.

This will become an error in seaborn v0.14.0; please update your code.

```
kdeplot(**{axis: a}, ax=ax, color=kde_color, **kde_kws)
```




```
fig, ax = plt.subplots(figsize=(20,10))
ax.bar(data.columns.map(str), data.mean().values)
ax.set_xlabel('Variables')
plt.xticks(rotation = 90)
plt.title('Mean of each variable of the dataset');
```




```
data_melted = pd.melt(data)
```

```
sns.set(rc={"figure.figsize":(20, 10)}) #width=3, #height=4  
sns.boxplot(x='variable', y='value', data=data_melted);
```




```
# Correlation between the release tahun and features
corr = data.corr()
fig, ax = plt.subplots(figsize=(10,10))
plt.title("Correlation matrix")
sns.heatmap(corr, square=True);
```



```
for t in column_names[1:6]:
    y = data[t].to_numpy()
    plt.plot(sorted(y), label=t, linewidth=1)

plt.legend(loc='upper center', bbox_to_anchor=(0.3, 1.03), ncol=3, fancybox=True, shadow=True)
```




```
X = data.iloc[:, 1:].to_numpy()
X = (X - X.min()) / (X.max() - X.min())


for i in range(1, 51):
    plt.plot(X[i], label='X' + str(i))

plt.xlabel("Feature Label")
plt.ylabel("Value Label")
plt.legend(loc='upper center', bbox_to_anchor=(0.8, 0.9), ncol=5, fancybox=True, shadow=True, fontsize=7)
```


 <matplotlib.legend.Legend at 0x7cd700d62f80>



```
fig, ax = plt.subplots(figsize=(25,7))
sns.barplot(x=corr['tahun'][1:].index, y=corr['tahun'][1:].values)
plt.title('Correlation to tahun');
```



✎ Splitting the data set

We split the data set into a training and a testing data set, before applying any pre-processing of the data, as it would otherwise put information from the testing set into the training set.

We follow the instruction given on the data set page on the UCI Machine Learning Repository and split the data set this way :

train: first 463,715 examples

test: last 51,630 examples

Which according to the website "avoids the 'producer effect' by making sure no song from a given artist ends up in both the train and test set."

```
data_train=data.iloc[:463715,:]
print(data_train.shape)
data_test=data.iloc[463715:,:]
print(data_test.shape)
```

```
(463715, 91)
(51416, 91)
```

```
data_train.describe()
```

```
(463715, 91)
(51416, 91)
```

	tahun	x1	x2	x3	x4	x5	x6	x7
count	463715.000000	463715.000000	463715.000000	463715.000000	463715.000000	463715.000000	463715.000000	463715.000000
mean	1998.386492	43.385407	1.253786	8.651491	1.130590	-6.513477	-9.566442	-2.383797
std	10.940319	6.079760	51.612880	35.264853	16.334058	22.854770	12.836177	14.580237
min	1922.000000	1.749000	-337.092500	-301.005060	-154.183580	-181.953370	-81.794290	-188.214000
25%	1994.000000	39.957510	-26.161450	-11.441550	-8.514270	-20.635820	-18.469000	-10.774800
50%	2002.000000	44.262120	8.361490	10.472720	-0.691060	-5.993610	-11.209400	-2.047330
75%	2006.000000	47.833710	36.136950	29.744940	8.756665	7.745720	-2.423955	6.516025
max	2011.000000	61.970140	384.065730	322.851430	289.527430	262.068870	119.815590	172.402680

8 rows × 91 columns

```
# create the scaler
ss = preprocessing.StandardScaler()


# create new dataframes to keep the non scaled ones
data_train_s=data_train.copy()
data_test_s=data_test.copy()

# apply the scaler to the dataframe subset
data_train_s.iloc[:,1:] = ss.fit_transform(data_train_s.iloc[:,1:])
data_test_s.iloc[:,1:] = ss.transform(data_test_s.iloc[:,1:])


means = pd.DataFrame(list(zip(data_train.columns, data_train.mean(), data_test.mean()))),
                      columns=['variable', 'train', 'test'])
means.drop(0, inplace=True)

means = pd.melt(means, id_vars="variable", var_name="dataframe", value_name="mean")

sns.catplot(x='variable', y='mean', hue='dataframe', data=means, kind='bar', height=8, aspect=1.7)
plt.xticks(rotation = 90)
plt.title('Mean of each feature in the training and testing dataframes before scaling');
```




```
means = pd.DataFrame(list(zip(data_train.columns, data_train_s.mean(), data_test_s.mean())),\n                        columns=['variable', 'train scaled', 'test scaled'])\nmeans.drop(0, inplace=True)\n\nmeans = pd.melt(means, id_vars="variable", var_name="dataframe", value_name="mean")\n\nsns.catplot(x='variable', y='mean', hue='dataframe', data=means, kind='bar', height=8, aspect=1.7)\nplt.xticks(rotation = 90)\nplt.title('Mean of each feature in the training and testing dataframes after scaling');
```



```
maxi = pd.DataFrame(list(zip(data_train.columns, data_train.max(), data_test.max()))),
                    columns=['variable', 'train', 'test'])
maxi.drop(0, inplace=True)

maxi = pd.melt(maxi, id_vars="variable", var_name="dataframe", value_name="maximum")


sns.catplot(x='variable', y='maximum', hue='dataframe', data=maxi, kind='bar', height=8, aspect=1.7)
plt.xticks(rotation = 90)
plt.title('Maximum of each feature in the training and testing dataframes before scaling');
```



```
maxi = pd.DataFrame(list(zip(data_train.columns, data_train_s.max(), data_test_s.max()),
                          columns=['variable', 'train scaled', 'test scaled']))
maxi.drop(0, inplace=True)

maxi = pd.melt(maxi, id_vars="variable", var_name="dataframe", value_name="maximum")

sns.catplot(x='variable', y='maximum', hue='dataframe', data=maxi, kind='bar', height=8, aspect=1.7)
plt.xticks(rotation = 90)
plt.title('Maximum of each feature in the training and testing dataframes after scaling');
```



```
# Downsampling by 'tahun'
min_samples = 1000
tahuns = data_train_s.tahun.unique()
sampled_dfs = [] # List to store each sampled DataFrame

for tahun in tahuns:
    if data_train_s[data_train_s.tahun == tahun].shape[0] > min_samples:
        sampled_dfs.append(data_train_s[data_train_s.tahun == tahun].sample(min_samples))
    else:
        sampled_dfs.append(data_train_s[data_train_s.tahun == tahun])

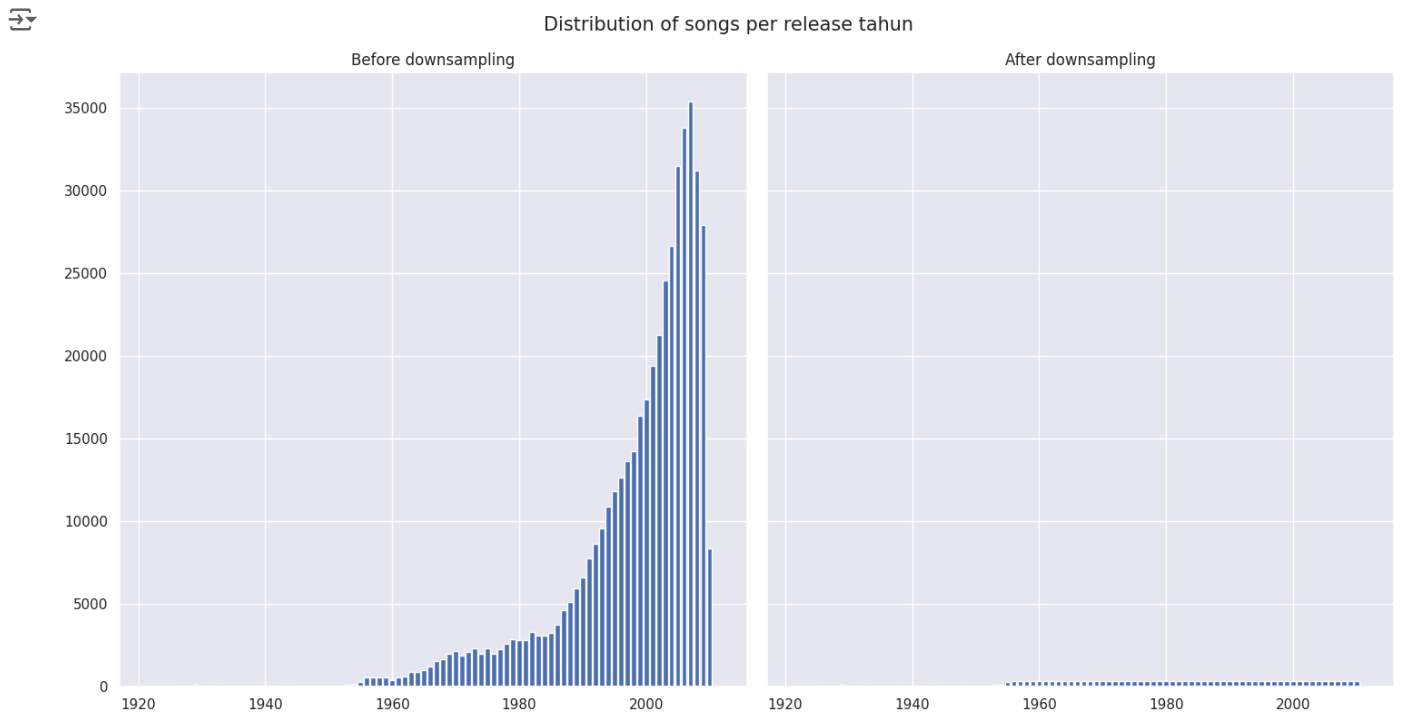
# Concatenate all sampled DataFrames at once
data_train_sampled = pd.concat(sampled_dfs, ignore_index=True)

fig, ax = plt.subplots(nrows=1, ncols=2, sharey=True, figsize=(15,8))
fig.suptitle('Distribution of songs per release tahun', fontsize=15)

ax[0].bar(data_train_s.tahun.value_counts().index, data_train_s.tahun.value_counts())
ax[0].set_title('Before downsampling')

ax[1].bar(data_train_sampled.tahun.value_counts().index, data_train_sampled.tahun.value_counts())
ax[1].set_title('After downsampling')

plt.tight_layout()
```



```
del data
```

```
data_train_s.shape
data_train_sampled.shape
```

```
(17909, 91)
```

```
from sklearn.ensemble import ExtraTreesClassifier
```

```
# we separate the target from the features
```

```
X_train = data_train_s.iloc[:,1:]
```

```
y_train = data_train_s.iloc[:,0]
```

```
model = ExtraTreesClassifier(n_estimators=10, max_depth=10, warm_start=True)
```

```
for i in range(10, 51, 10): # Increasing trees gradually
```


```
    model.n_estimators = i
```

```
    model.fit(X_train, y_train)
```

```
# graph of the 10 most important features
```

```
feat_importances = pd.Series(model.feature_importances_, index=X_train.columns)
```


```
feat_importances.nlargest(10).plot(kind='barh')
```




```
# graph of the 20 most important features
feat_importances = pd.Series(model.feature_importances_, index=X_train.columns)
feat_importances.nlargest(20).plot(kind='barh')
```


 <Axes: >

```
names10=['tahun']  
names10.extend(list(feats_importances.nlargest(10).index.sort_values()))  
names10
```

 ['tahun', 'x1', 'x14', 'x2', 'x20', 'x22', 'x3', 'x4', 'x6', 'x63', 'x7']

```
names20=['tahun']  
names20.extend(list(feats_importances.nlargest(20).index.sort_values()))  
names20
```

 ['tahun',
 'x1',
 'x14']